

Lab 10

Attila Lazar

09.12.2020

data

We load the dataset, and define a train and a test set.

```
data(Auto, package="ISLR")
```

```
set.seed(123)
```

```
n <- nrow(Auto)
```

```
train <- sample(1:n, round(n*2/3))
```

```
test <- (1:n) [-train]
```

```
str(Auto)
```

```
## 'data.frame':   392 obs. of  9 variables:
##  $ mpg          : num  18 15 18 16 17 15 14 14 15 ...
##  $ cylinders    : num  8 8 8 8 8 8 8 8 8 ...
##  $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
##  $ horsepower   : num  130 165 150 150 140 198 220 215 225 190 ...
##  $ weight       : num  3504 3693 3436 3433 3449 ...
##  $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
##  $ year         : num  70 70 70 70 70 70 70 70 70 70 ...
##  $ origin       : num  1 1 1 1 1 1 1 1 1 ...
##  $ name         : Factor w/ 304 levels "amc ambassador brougham",...: 49 36 231 14 161 141 54 223 241 ...
```

1)

We train a linear model using *lm* and *natural cubic splines*. Since cylinders and origin are categorical variables, they enter the model linearly.

```
library(splines)
```

```
model1 <- lm(mpg ~ ns(displacement, 4) + ns(horsepower, 4) + ns(acceleration, 4) + ns(weight, 4) + origin
```

a)

we interpret the model using *summary*

significantly contributing variables are *horsepower*, *year*, *weight* and *acceleration*

```
summary(model1)
```

```
##
```

```
## Call:
```

```
## lm(formula = mpg ~ ns(displacement, 4) + ns(horsepower, 4) +
```

```
##      ns(acceleration, 4) + ns(weight, 4) + origin + ns(year, 4) +
```

```
##      cylinders, data = Auto, subset = train)
```


```
##
```

```
## Residuals:
```

```
##      Min      1Q  Median      3Q      Max
## -7.8917 -1.4840  0.1298  1.4435  7.4804
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      40.1461      3.8229  10.501 < 2e-16 ***
## ns(displacement, 4)1 -2.3655      2.2176  -1.067 0.287185
## ns(displacement, 4)2 -3.1748      2.9829  -1.064 0.288268
## ns(displacement, 4)3 -2.3071      4.2553  -0.542 0.588212
## ns(displacement, 4)4 -5.3432      3.4005  -1.571 0.117445
## ns(horsepower, 4)1  -8.1527      1.9244  -4.237 3.25e-05 ***
## ns(horsepower, 4)2 -12.6926      2.4476  -5.186 4.60e-07 ***
## ns(horsepower, 4)3 -22.5621      4.2721  -5.281 2.89e-07 ***
## ns(horsepower, 4)4 -11.9723      2.6830  -4.462 1.25e-05 ***
## ns(acceleration, 4)1 -6.2436      2.7551  -2.266 0.024340 *
## ns(acceleration, 4)2 -5.5801      1.9391  -2.878 0.004370 **
## ns(acceleration, 4)3 -11.2939      5.7682  -1.958 0.051403 .
## ns(acceleration, 4)4 -4.0390      2.5383  -1.591 0.112889
## ns(weight, 4)1      -7.9846      2.1543  -3.706 0.000261 ***
## ns(weight, 4)2      -7.4575      2.5580  -2.915 0.003892 **
## ns(weight, 4)3     -11.4807      4.4505  -2.580 0.010491 *
## ns(weight, 4)4      -6.2676      2.9929  -2.094 0.037304 *
## origin              0.5120      0.3126   1.638 0.102748
## ns(year, 4)1        -0.3904      0.8922  -0.438 0.662109
## ns(year, 4)2         6.5663      0.8607   7.629 5.67e-13 ***
## ns(year, 4)3         6.1018      1.6303   3.743 0.000228 ***
## ns(year, 4)4         7.9807      0.7106  11.232 < 2e-16 ***
## cylinders           0.4917      0.4438   1.108 0.269070
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.631 on 238 degrees of freedom
## Multiple R-squared:  0.9008, Adjusted R-squared:  0.8916
## F-statistic: 98.25 on 22 and 238 DF,  p-value: < 2.2e-16
```

We see on the validation plot that the model predicts the test data quite well

```
pred1 <- predict(model1, Auto[test,])
plot(Auto[test,'mpg'], pred1, xlab='y' ,ylab='y-hat', main="validation")
abline(c(0,1))
```



with RMSE of

```
sqrt(mean((Auto[test, 'mpg'] - predict(model1, Auto[test,]))^2))
```

```
## [1] 2.996651
```

b)

Now we use stepwise reduction of the model.

```
model2 <- step(lm(mpg~ns(displacement, 4) + ns(horsepower, 4) + ns(acceleration, 4) + ns(weight, 4) + origin, data = Auto, subset = train))
summary(model2)
```

```
##
## Call:
## lm(formula = mpg ~ ns(horsepower, 4) + ns(acceleration, 4) +
##      ns(weight, 4) + origin + ns(year, 4), data = Auto, subset = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-9.0631	-1.3062	0.1386	1.4229	7.5354

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	39.8651	2.9659	13.441	< 2e-16 ***
ns(horsepower, 4)1	-7.9741	1.8850	-4.230	3.31e-05 ***
ns(horsepower, 4)2	-11.0185	2.2895	-4.813	2.62e-06 ***
ns(horsepower, 4)3	-20.3118	4.1708	-4.870	2.01e-06 ***
ns(horsepower, 4)4	-11.3916	2.6121	-4.361	1.91e-05 ***

```
## ns(acceleration, 4)1 -4.9883      2.3906 -2.087 0.037968 *
## ns(acceleration, 4)2 -4.8724      1.7464 -2.790 0.005689 **
## ns(acceleration, 4)3 -8.7794      5.1606 -1.701 0.090176 .
## ns(acceleration, 4)4 -2.4111      2.3575 -1.023 0.307463
## ns(weight, 4)1 -8.4890      1.7382 -4.884 1.89e-06 ***
## ns(weight, 4)2 -9.5785      1.8445 -5.193 4.37e-07 ***
## ns(weight, 4)3 -12.4559      3.7075 -3.360 0.000906 ***
## ns(weight, 4)4 -9.0271      2.1770 -4.147 4.67e-05 ***
## origin          0.7628      0.2743  2.781 0.005851 **
## ns(year, 4)1 -0.1018      0.8745 -0.116 0.907416
## ns(year, 4)2  6.5648      0.8400  7.816 1.65e-13 ***
## ns(year, 4)3  6.1643      1.5999  3.853 0.000149 ***
## ns(year, 4)4  8.1951      0.6958 11.777 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.639 on 243 degrees of freedom
## Multiple R-squared:  0.8981, Adjusted R-squared:  0.891
## F-statistic: 126 on 17 and 243 DF, p-value: < 2.2e-16
```

This eliminates the variable *displacement* and gives us a slightly better RMSE

```
sqrt(mean((Auto[test, 'mpg'] - predict(model2, Auto[test,]))^2))
```

```
## [1] 2.961918
```

c)

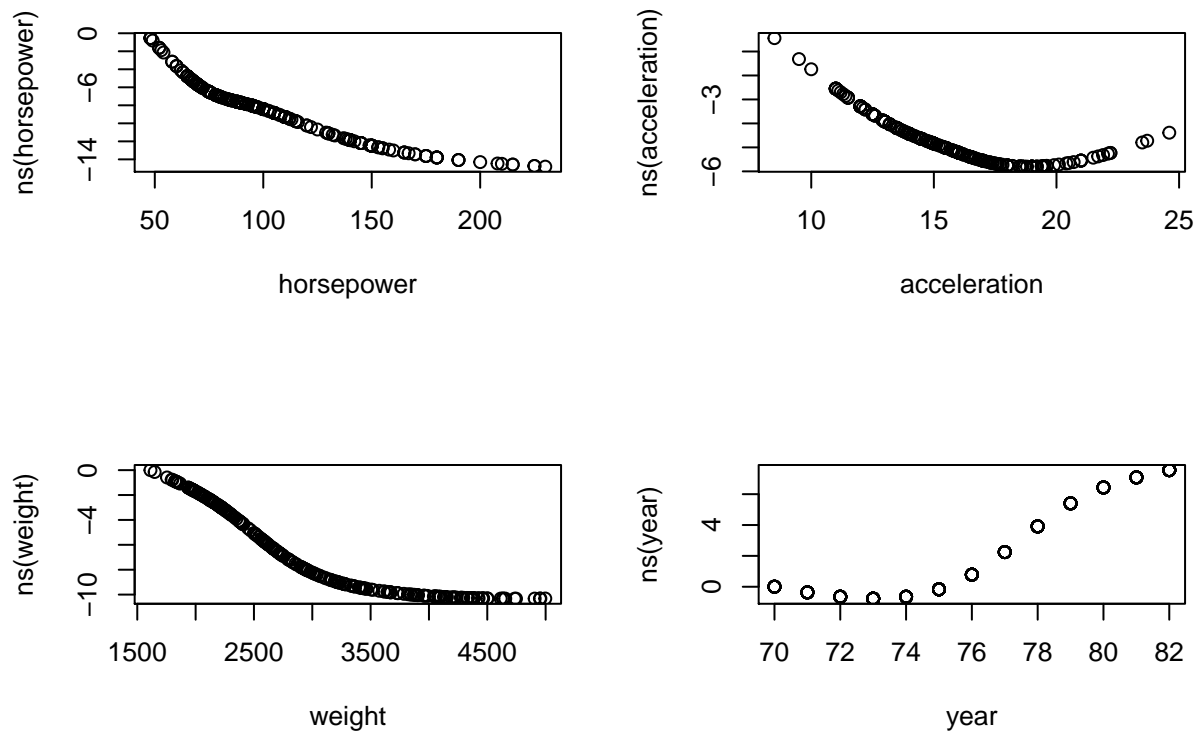
The model consists of the Intercept, the coefficients for each spline, and a coefficient for linearly modeled variables.

```
model2$coefficients
```

```
##      (Intercept) ns(horsepower, 4)1 ns(horsepower, 4)2
##      39.8651495      -7.9740653      -11.0185477
## ns(horsepower, 4)3 ns(horsepower, 4)4 ns(acceleration, 4)1
##      -20.3117557      -11.3915600      -4.9882998
## ns(acceleration, 4)2 ns(acceleration, 4)3 ns(acceleration, 4)4
##      -4.8723536      -8.7794202      -2.4110593
## ns(weight, 4)1 ns(weight, 4)2 ns(weight, 4)3
##      -8.4889941      -9.5785026      -12.4559214
## ns(weight, 4)4 origin ns(year, 4)1
##      -9.0270585      0.7627804      -0.1018131
## ns(year, 4)2 ns(year, 4)3 ns(year, 4)4
##      6.5648195      6.1643067      8.1951039
```

We plot the calculated value of the splines in the model against the original variable

```
par(mfrow=c(2,2))
plot(Auto$horsepower[train], model2$model$`ns(horsepower, 4)` %>% model2$coefficients[2:5], xlab='horsepower', ylab='ns(horsepower, 4)')
plot(Auto$acceleration[train], model2$model$`ns(acceleration, 4)` %>% model2$coefficients[6:9], xlab='acceleration', ylab='ns(acceleration, 4)')
plot(Auto$weight[train], model2$model$`ns(weight, 4)` %>% model2$coefficients[10:13], xlab='weight', ylab='ns(weight, 4)')
plot(Auto$year[train], model2$model$`ns(year, 4)` %>% model2$coefficients[15:18], xlab='year', ylab='ns(year, 4)')
```



In these plots we see how the variable enters the model. For *horsepower* and *weight* we see a near linear, negative trend which is expected.

Interestingly *acceleration* over 20 positively affects *mpg* reversing the trend. This may be because there are only few datapoints which may affect the model

Lastly *year* affects *mpg* negatively until 73, after that *year* strongly increases *mpg*. This may be attributed to the 1973 oil crisis.

2)

a)

We use *gam* to compute *Generalized Additive Models*. As in Ex1, we do not construct splines for *origin* and *cylinders*, since these are categorical variables.

```
library(mgcv)
```

```
## Loading required package: nlme
```

```
## This is mgcv 1.8-24. For overview type 'help("mgcv-package")'.
```

```
model3 <- gam(mpg~s(displacement) + s(horsepower) + s(acceleration) + s(weight) + origin + s(year) + cy
```

b)

```
summary(model3)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## mpg ~ s(displacement) + s(horsepower) + s(acceleration) + s(weight) +
##       origin + s(year) + cylinders
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  21.4749    1.9928  10.776  <2e-16 ***
## origin        0.6412    0.2845   2.254  0.0251 *
## cylinders     0.1965    0.3611   0.544  0.5868
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F  p-value
## s(displacement) 1.315  1.543  0.549 0.385032
## s(horsepower)   2.863  3.648  5.946 0.000329 ***
## s(acceleration) 2.144  2.756  2.078 0.083199 .
## s(weight)       2.376  3.037 10.204 2.23e-06 ***
## s(year)         8.529  8.931 37.053 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.899   Deviance explained = 90.6%
## GCV = 7.0162   Scale est. = 6.4725     n = 261
```

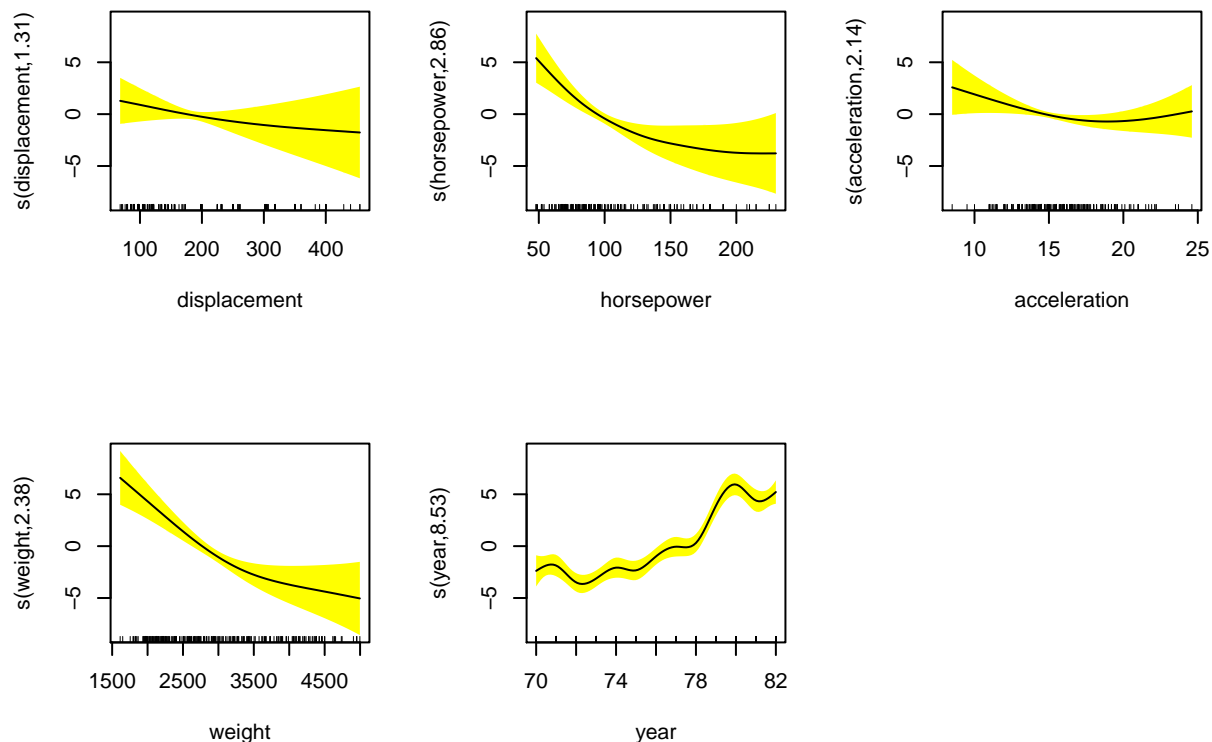
We see that the model attributes no significance to *displacement* and only little significance to *acceleration*. Also we see that the smooth function for *displacement* is near linear with $\text{edf} = 1.3$. In contrast it is quite complex for *year* with $\text{edf} = 8.5$.

c)

We plot the smmoth functions. We can see how the variable enters the model and how it affects the predicted variable.

The smooth function for *year* seems to be too complex which might lead to overfitting.

```
plot(model13, page=1,shade=TRUE,shade.col = "yellow")
```



d)

```
sqrt(mean((Auto[test, 'mpg'] - predict(model3, Auto[test,]))^2))
```

```
## [1] 2.90203
```

e)

first we try to enhance our model by manually restricting the choice of k value for *year*. We see a good improvement of the RMSE and also the complexity of the smooth function is reduced.

```
model5 <- gam(mpg~s(displacement) + s(horsepower) + s(acceleration) + s(weight) + origin + s(year, k=3)
#plot(model5, page=1,shade=TRUE,shade.col = "yellow")
summary(model5)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## mpg ~ s(displacement) + s(horsepower) + s(acceleration) + s(weight) +
##       origin + s(year, k = 3) + cylinders
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  20.7398     2.0464   10.13  <2e-16 ***
## origin        0.6953     0.2959    2.35   0.0196 *
```

```
## cylinders      0.3154      0.3712      0.85   0.3964
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(displacement) 1.187  1.331   0.628 0.363472
## s(horsepower)    2.837  3.613   6.753 0.000103 ***
## s(acceleration) 2.254  2.899   2.626 0.041460 *
## s(weight)        2.571  3.282  10.035 1.83e-06 ***
## s(year)          1.962  1.998 124.006 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.885   Deviance explained = 89.1%
## GCV = 7.7478   Scale est. = 7.3378      n = 261
sqrt(mean((Auto[test, 'mpg'] - predict(model5, Auto[test,]))^2))

## [1] 2.838914
```

Next we try the option bs=ts. This results in similar model than our original smooth model.

```
model7 <- gam(mpg~s(displacement,bs='ts') + s(horsepower,bs='ts') + s(acceleration,bs='ts') + s(weight,bs='ts') + s(year,bs='ts') + cylinders,
#plot(model7, page=1,shade=TRUE,shade.col = "yellow")
summary(model7)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## mpg ~ s(displacement, bs = "ts") + s(horsepower, bs = "ts") +
##       s(acceleration, bs = "ts") + s(weight, bs = "ts") + origin +
##       s(year, bs = "ts") + cylinders
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 22.81788    1.63261  13.976 < 2e-16 ***
## origin       0.73922    0.27140   2.724 0.00693 **
## cylinders   -0.07633    0.27534  -0.277 0.78185
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(displacement) 0.07847     9  0.008  0.3226
## s(horsepower)    3.12305     9  2.469 7.68e-06 ***
## s(acceleration) 2.09404     9  0.674  0.0302 *
## s(weight)        2.52872     9  5.199 3.62e-12 ***
## s(year)          8.52977     9 36.514 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.899   Deviance explained = 90.6%
## GCV = 7.001   Scale est. = 6.4818      n = 261
```



```
sqrt(mean((Auto[test, 'mpg'] - predict(model7, Auto[test,]))^2))
```

```
## [1] 2.907251
```

We also try the option `bs='cr'`. The smooth function complexities are reduced slightly and we improve the RMSE

```
model8 <- gam(mpg~s(displacement,bs='cr') + s(horsepower,bs='cr') + s(acceleration,bs='cr') + s(weight,bs='cr') + s(year,bs='cr') + s(cylinders,bs='cr'),
#plot(model8, page=1,shade=TRUE,shade.col = "yellow")
summary(model8)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## mpg ~ s(displacement, bs = "cr") + s(horsepower, bs = "cr") +
##       s(acceleration, bs = "cr") + s(weight, bs = "cr") + origin +
##       s(year, bs = "cr") + cylinders
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  21.7625     2.1100  10.314   <2e-16 ***
## origin        0.5881     0.2933   2.005   0.0461 *
## cylinders     0.1591     0.3824   0.416   0.6777
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F  p-value
## s(displacement) 1.427  1.728  1.008 0.436540
## s(horsepower)   3.276  4.132  5.192 0.000423 ***
## s(acceleration) 2.749  3.515  2.823 0.037482 *
## s(weight)       3.399  4.297  7.328 8.98e-06 ***
## s(year)         5.099  6.159 47.091 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.894   Deviance explained = 90.1%
## -REML = 624.76   Scale est. = 6.79        n = 261
```

```
sqrt(mean((Auto[test, 'mpg'] - predict(model8, Auto[test,]))^2))
```

```
## [1] 2.856505
```

Now set the option `select = TRUE`. We arrive at our best RMSE.

```
model9 <- gam(mpg~s(displacement,bs='cr') + s(horsepower,bs='cr') + s(acceleration,bs='cr') + s(weight,bs='cr') + s(year,bs='cr') + s(cylinders,bs='cr'),
#plot(model9, page=1,shade=TRUE,shade.col = "yellow")
summary(model9)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## mpg ~ s(displacement, bs = "cr") + s(horsepower, bs = "cr") +
##       s(acceleration, bs = "cr") + s(weight, bs = "cr") + s(year, bs = "cr") + s(cylinders, bs = "cr")
```

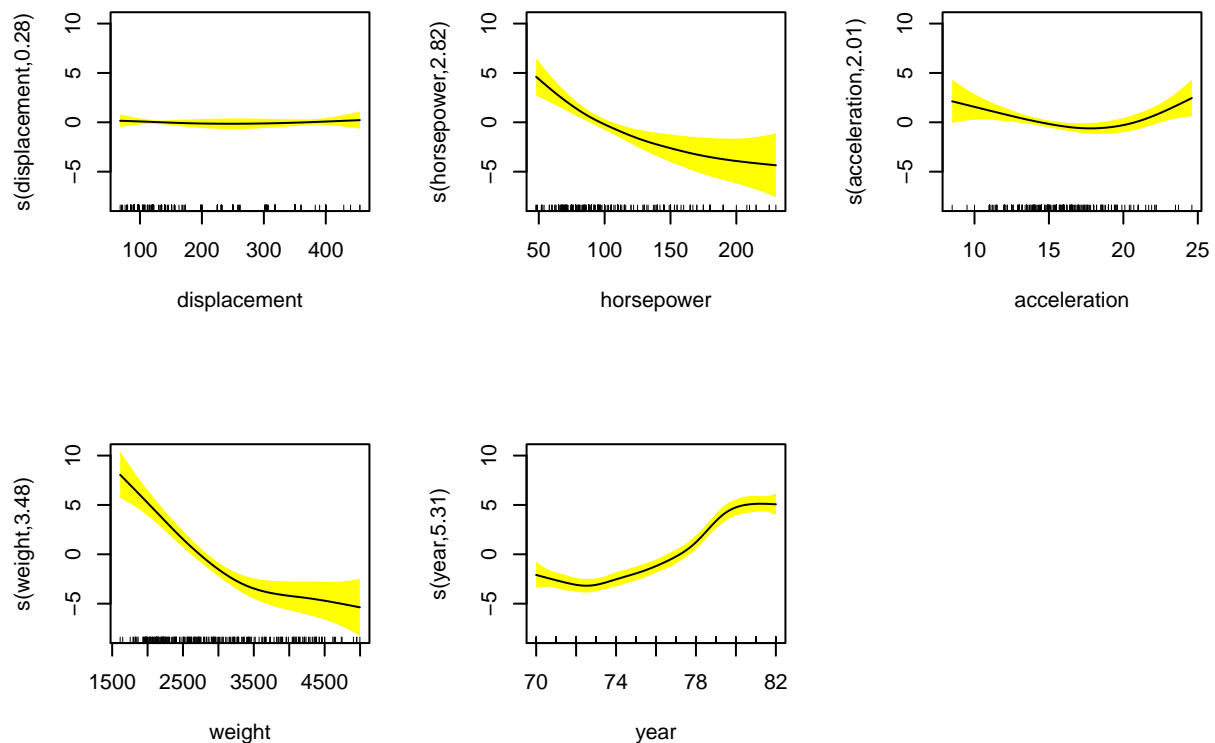
```
##      s(acceleration, bs = "cr") + s(weight, bs = "cr") + origin +
##      s(year, bs = "cr") + cylinders
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 22.96481    1.69648  13.537  <2e-16 ***
## origin      0.66607     0.27716   2.403   0.017 *
## cylinders   -0.08243     0.28540  -0.289   0.773
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F  p-value
## s(displacement) 0.281     9  0.037  0.23046
## s(horsepower)   2.821     9  2.943 1.35e-07 ***
## s(acceleration) 2.010     9  0.961  0.00291 **
## s(weight)       3.476     9  7.708 < 2e-16 ***
## s(year)         5.314     9 33.370 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.894   Deviance explained =  90%
## -REML = 640.45   Scale est. = 6.7912     n = 261
```

```
sqrt(mean((Auto[test, 'mpg'] - predict(model9, Auto[test,]))^2))
```

```
## [1] 2.83679
```

We plot the smooth functions.

```
plot(model9, page=1,shade=TRUE,shade.col = "yellow")
```



```
plot(Auto[test,'mpg'], predict(model9, Auto[test,]), xlab='y' ,ylab='y-hat', main="validation")  
abline(c(0,1))
```

